

Recurrent Neural Networks

Girish Varma

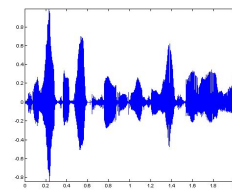
IIIT Hyderabad

<http://bit.ly/2ubR4iC>

Sequence to Sequence Problems

Convert x_1, x_2, \dots, x_s to y_1, y_2, \dots, y_t

Speech to Text: speech signal $x(i)$ to text $y(i)$



→ Hello

Language Translation: text in language 1 $x(i)$ to text in language 2 $y(i)$



बुश की 77 करोड़
डॉलर की मदद की
पेशकश

अमरीकी राष्ट्रपति बुश ने दुनियाभर में
खाद्य पदार्थों की बढ़ती कीमतों के असर
को कम करने के लिए 77 करोड़ डॉलर
की खाद्य सहायता की पेशकश की है.
+ खाद्यान्न संकट टास्क फ़ोर्स का गठन



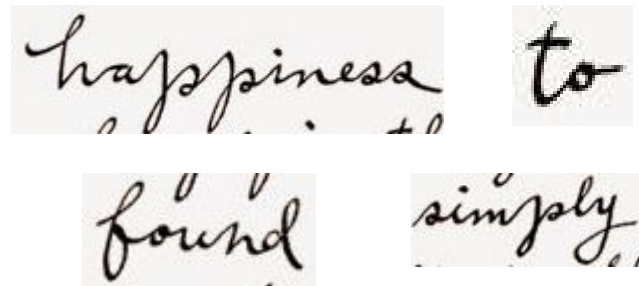
Bush offered the
help of a \$ 77
million

U.S. President George Bush has
in the world of food products to
reduce the impact of rising prices
for 77 million dollars in food aid
offered.

Sequence to Sequence Problems

Convert x_1, x_2, \dots, x_s to y_1, y_2, \dots, y_t

OCR : Convert these images to unicode strings



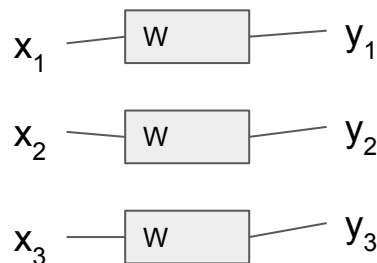
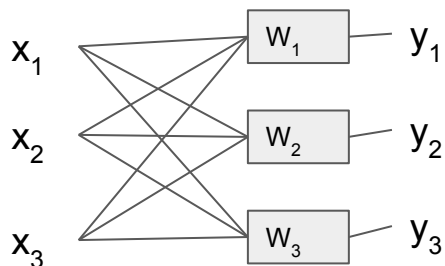
Time series data : Given a sequence predict the next element

- Predict stock prices
- Weather prediction



First Attempt at a Model

Convert x_1, x_2, x_3 to y_1, y_2, y_3



1. Variable length sequences.
2. Each network will have too many inputs.
3. Using a different Network each time, results in explosion of parameters.
Need information about previous inputs/outputs for predicting the next.

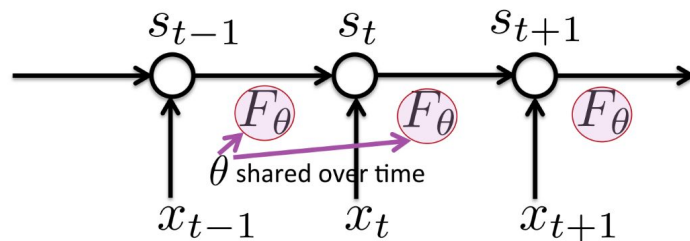
Recurrent Neural Networks

Convert x_1, x_2, \dots, x_s to o_1, o_2, \dots, o_t

Summarise input sequence in a fixed length state vector with recursive updates.

$$s_t = F_{\theta}(s_{t-1}, x_t)$$

$$o_t = G_{\theta'}(s_t)$$

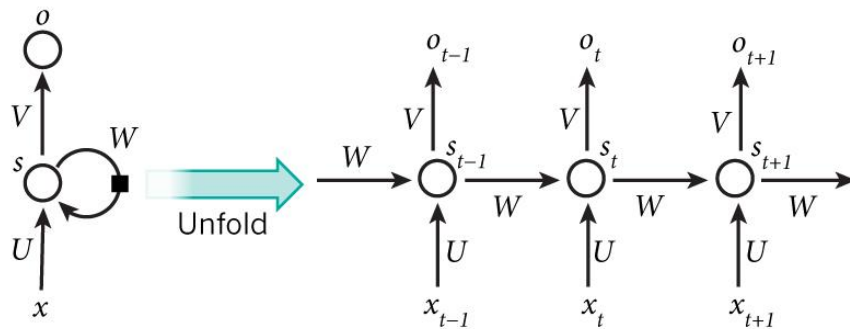


Generalizes to lengths not seen during training.

RNNs

$$\begin{aligned}s_t &= F_{U,W,b}(s_{t-1}, x_t) \\ &= \tanh(Ux_t + Ws_{t-1} + b)\end{aligned}$$

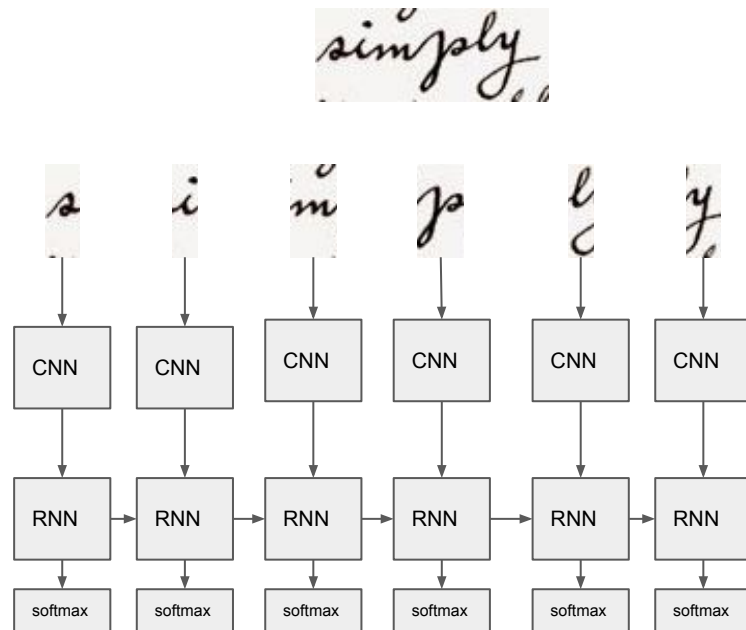
$$\begin{aligned}o_t &= G_V(s_t) \\ &= Vs_t\end{aligned}$$



Source: Nature

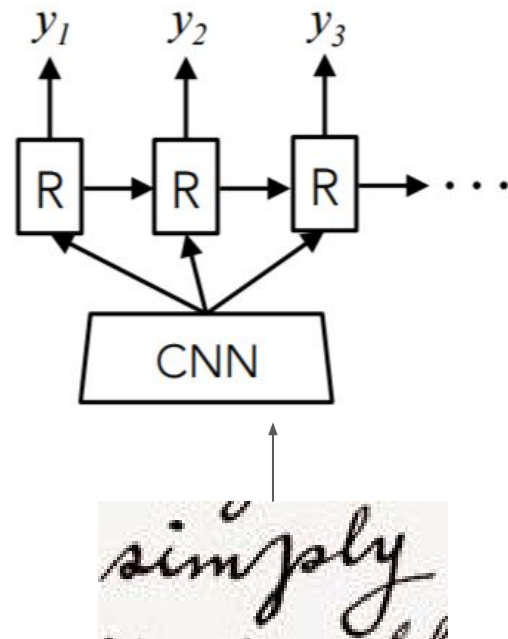
Case Study : Scene Text with char segmentation

1. Segment the image into characters using some algorithm.
2. Extract CNN features
3. Pass CNN feature sequence through RNN
4. Pass RNN output through softmax to get alphabet probabilities
5. Loss function: $\sum_i \text{Error}(y_i, y_i^{\text{correct}})$



Case Study : Scenetxt without char segmentation

1. CNN for image feature extraction.
2. Duplicate CNN features into a sequence of Maxlen.
3. RNN followed by softmax gives probabilities of alphabets of length Maxlen.



Backpropagation in RNNs

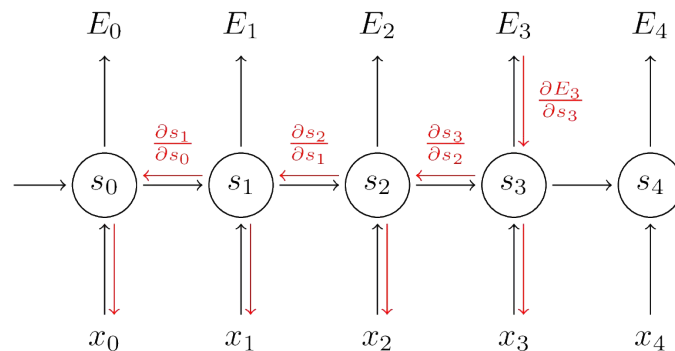
$$h_t = Wf(h_{t-1}) + Vx_t$$

- Total error is the sum of each error at time steps t

$$\frac{\partial E}{\partial W} = \sum_{t=1}^T \frac{\partial E_t}{\partial W}$$

- Hardcore chain rule application:

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial h_k}{\partial W}$$



The vanishing gradient problem - Details

- Similar to backprop but less efficient formulation
- Useful for analysis we'll look at:

$$\frac{\partial E_t}{\partial W} = \sum_{k=1}^t \frac{\partial E_t}{\partial y_t} \frac{\partial y_t}{\partial h_t} \boxed{\frac{\partial h_t}{\partial h_k}} \frac{\partial h_k}{\partial W}$$

- Remember: $h_t = Wf(h_{t-1}) + Vx_t$
- More chain rule, remember:

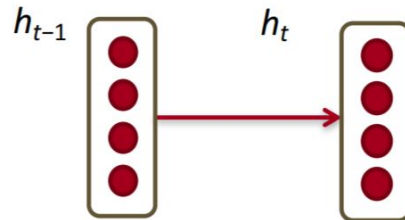
$$\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}}$$

- Each partial is a Jacobian:

$$\frac{d\mathbf{f}}{d\mathbf{x}} = \left[\frac{\partial \mathbf{f}}{\partial x_1} \quad \cdots \quad \frac{\partial \mathbf{f}}{\partial x_n} \right] = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

The vanishing gradient problem - Details

- From previous slide: $\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}}$



- Remember: $h_t = Wf(h_{t-1}) + Vx_t$

- To compute Jacobian, derive each element of matrix: $\frac{\partial h_{j,m}}{\partial h_{j-1,n}}$

$$\frac{\partial h_t}{\partial h_k} = \prod_{j=k+1}^t \frac{\partial h_j}{\partial h_{j-1}} = \prod_{j=k+1}^t W^T \text{diag}[f'(h_{j-1})]$$

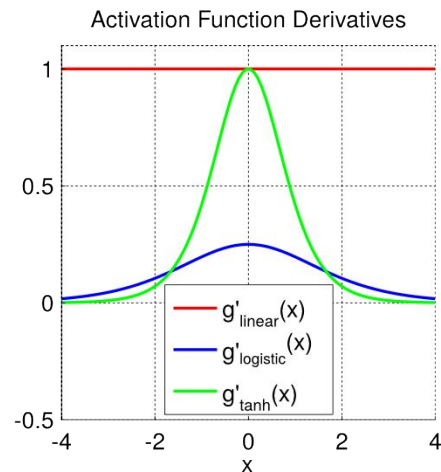
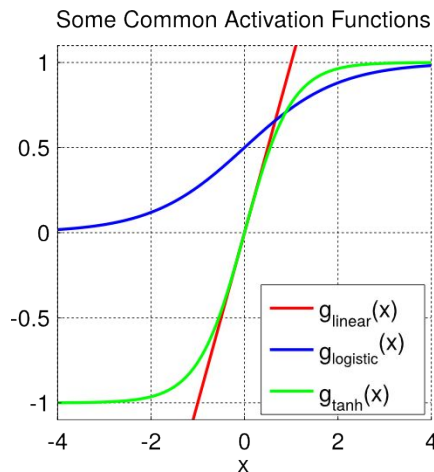
- Where: $\text{diag}(z) = \begin{pmatrix} z_1 & & & & \\ & z_2 & & 0 & \\ & & \ddots & & \\ & 0 & & z_{n-1} & \\ & & & & z_n \end{pmatrix}$

Check at home
that you understand
the diag matrix
formulation

Vanishing gradients

- If $t - k = d$, gradient will have a term of W^d .
- If $\|W\|$ is small, all long range dependencies will be lost.
- Tanh is used as activation function to prevent this problem from worsening

$$s_t = \tanh(Ux_t + Ws_{t-1} + b)$$



Gated Recurrent Units : GRUs

- More complex hidden unit computation in recurrence!
- Gated Recurrent Units (GRU) introduced by Cho et al. 2014 (see reading list)
- Main ideas:
 - keep around memories to capture long distance dependencies
 - allow error messages to flow at different strengths depending on the inputs

GRUs

- Standard RNN computes hidden layer at next time step directly:

$$h_t = f \left(W^{(hh)} h_{t-1} + W^{(hx)} x_t \right)$$

- GRU first computes an update **gate** (another layer) based on current input word vector and hidden state

$$z_t = \sigma \left(W^{(z)} x_t + U^{(z)} h_{t-1} \right)$$

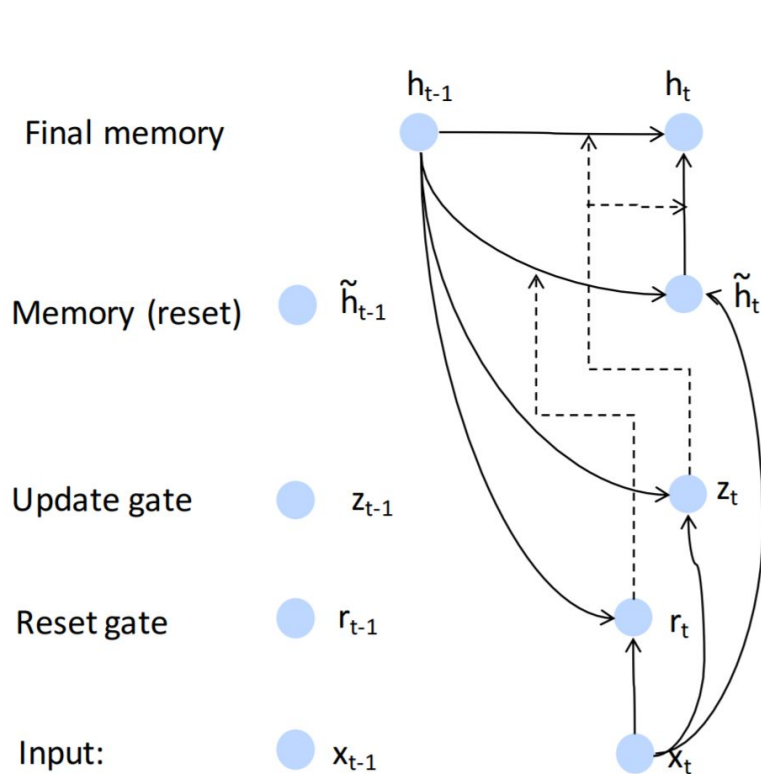
- Compute reset gate similarly but with different weights

$$r_t = \sigma \left(W^{(r)} x_t + U^{(r)} h_{t-1} \right)$$

GRUs

- Update gate $z_t = \sigma \left(W^{(z)} x_t + U^{(z)} h_{t-1} \right)$
- Reset gate $r_t = \sigma \left(W^{(r)} x_t + U^{(r)} h_{t-1} \right)$
- New memory content: $\tilde{h}_t = \tanh (W x_t + r_t \circ U h_{t-1})$
If reset gate unit is ~ 0 , then this ignores previous memory and only stores the new word information
- Final memory at time step combines current and previous time steps: $h_t = z_t \circ h_{t-1} + (1 - z_t) \circ \tilde{h}_t$

Attempt at a clean illustration



$$z_t = \sigma \left(W^{(z)} x_t + U^{(z)} h_{t-1} \right)$$

$$r_t = \sigma \left(W^{(r)} x_t + U^{(r)} h_{t-1} \right)$$

$$\tilde{h}_t = \tanh \left(W x_t + r_t \circ U h_{t-1} \right)$$

$$h_t = z_t \circ h_{t-1} + (1 - z_t) \circ \tilde{h}_t$$

GRU intuition

- If reset is close to 0, ignore previous hidden state
→ Allows model to drop information that is irrelevant in the future

$$\begin{aligned}z_t &= \sigma \left(W^{(z)} x_t + U^{(z)} h_{t-1} \right) \\r_t &= \sigma \left(W^{(r)} x_t + U^{(r)} h_{t-1} \right) \\\tilde{h}_t &= \tanh \left(W x_t + r_t \circ U h_{t-1} \right) \\h_t &= z_t \circ h_{t-1} + (1 - z_t) \circ \tilde{h}_t\end{aligned}$$

- Update gate z controls how much of past state should matter now.
 - If z close to 1, then we can copy information in that unit through many time steps! **Less vanishing gradient!**
- Units with short-term dependencies often have reset gates very active

GRU intuition

- Units with long term dependencies have active update gates z

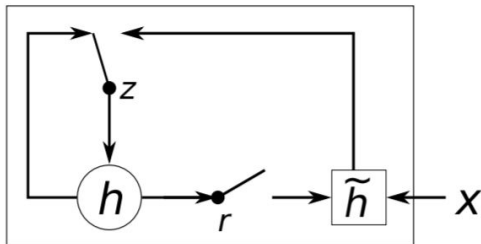
$$z_t = \sigma \left(W^{(z)} x_t + U^{(z)} h_{t-1} \right)$$

$$r_t = \sigma \left(W^{(r)} x_t + U^{(r)} h_{t-1} \right)$$

$$\tilde{h}_t = \tanh \left(W x_t + r_t \circ U h_{t-1} \right)$$

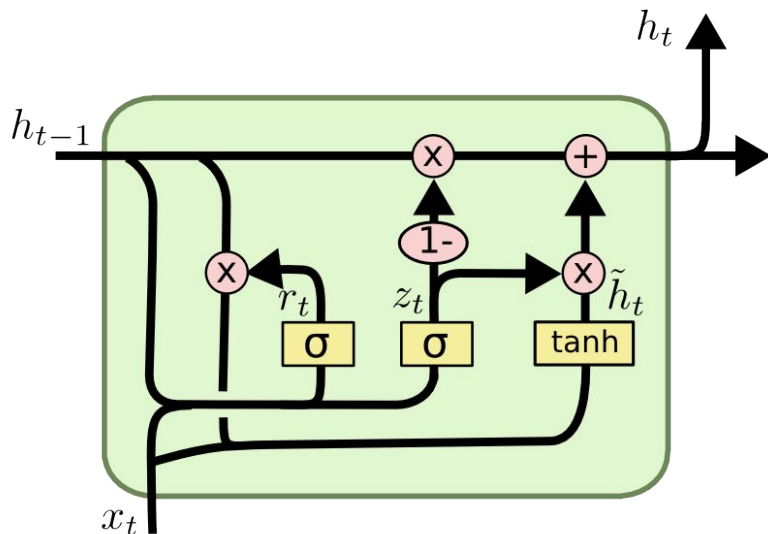
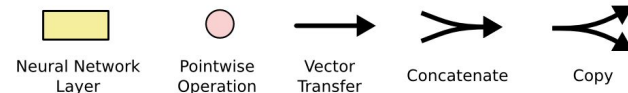
$$h_t = z_t \circ h_{t-1} + (1 - z_t) \circ \tilde{h}_t$$

- Illustration:



- Derivative of $\frac{\partial}{\partial x_1} x_1 x_2$? \rightarrow rest is same chain rule, but implement with **modularization** or automatic differentiation

GRU : Gated Recurrent Units



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

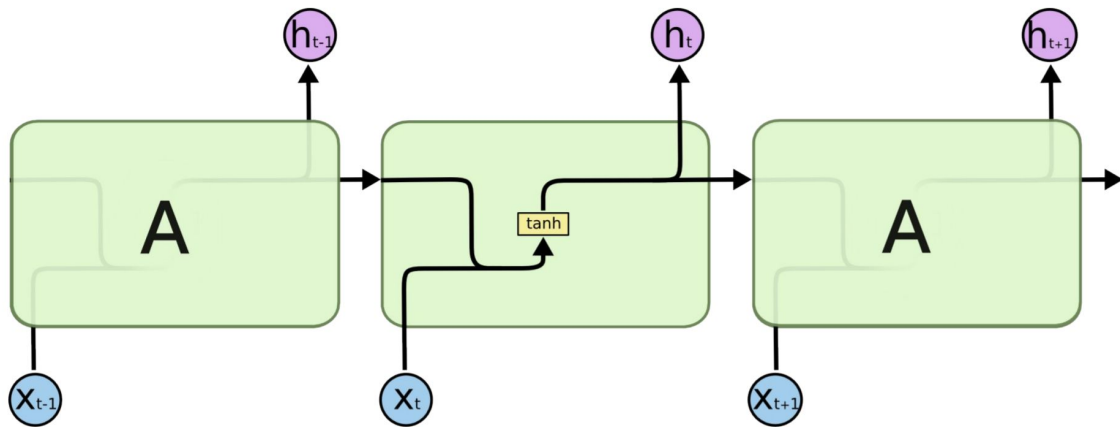
$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

LSTMs : Long Short Term Memory

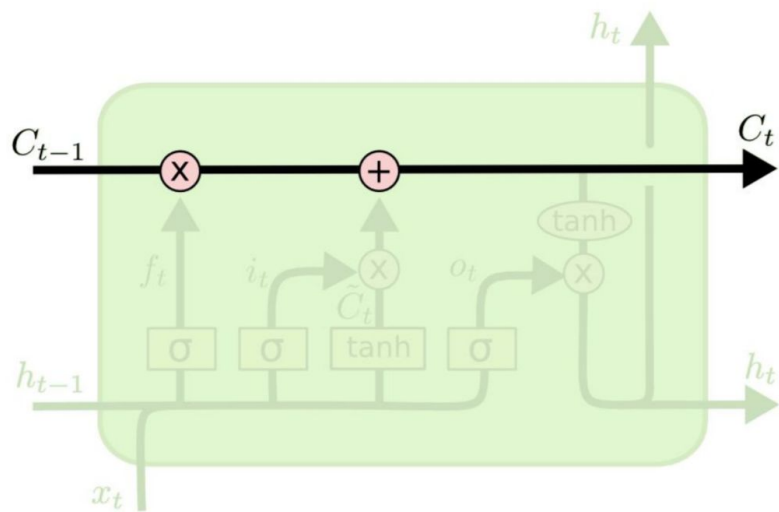
Hochreiter & Schmidhuber (1997) solved the problem of getting an RNN to remember things for a long time (like hundreds of time steps).

Based on a standard RNN whose neuron activates with *tanh*...



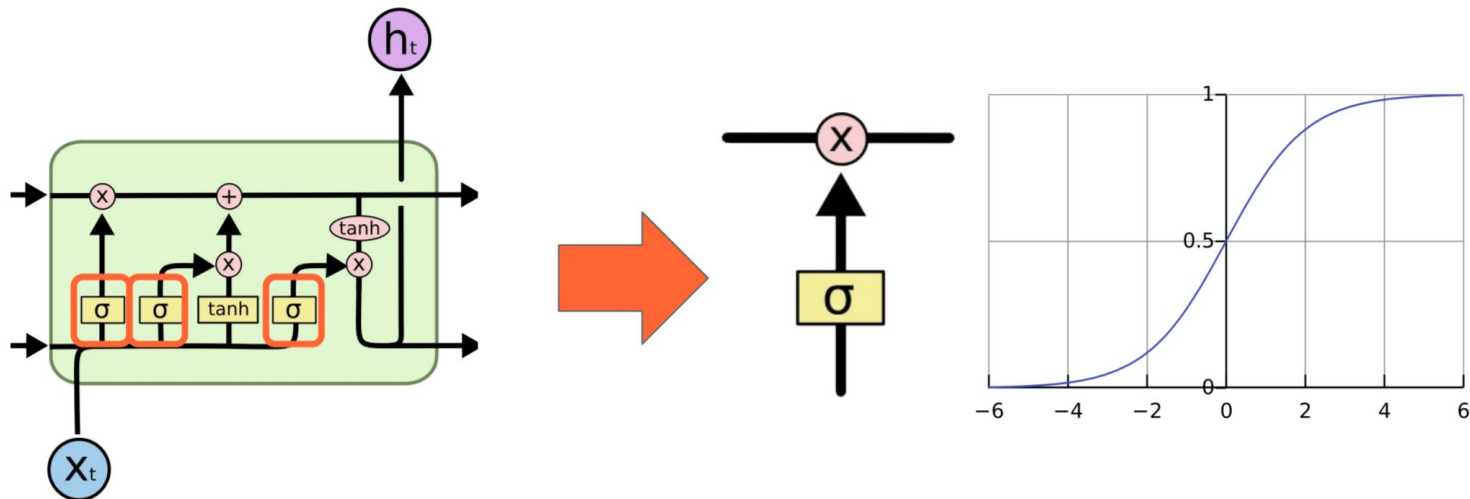
LSTM

C_t is the cell state, which flows through the entire chain...

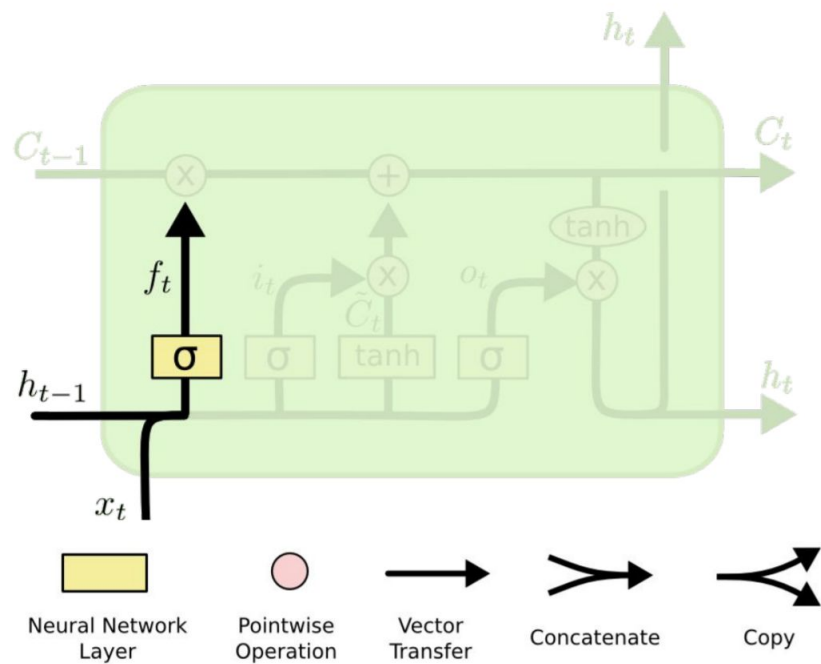


LSTM

Three **gates** are governed by *sigmoid* units (btw $[0,1]$) define the control of in & out information..



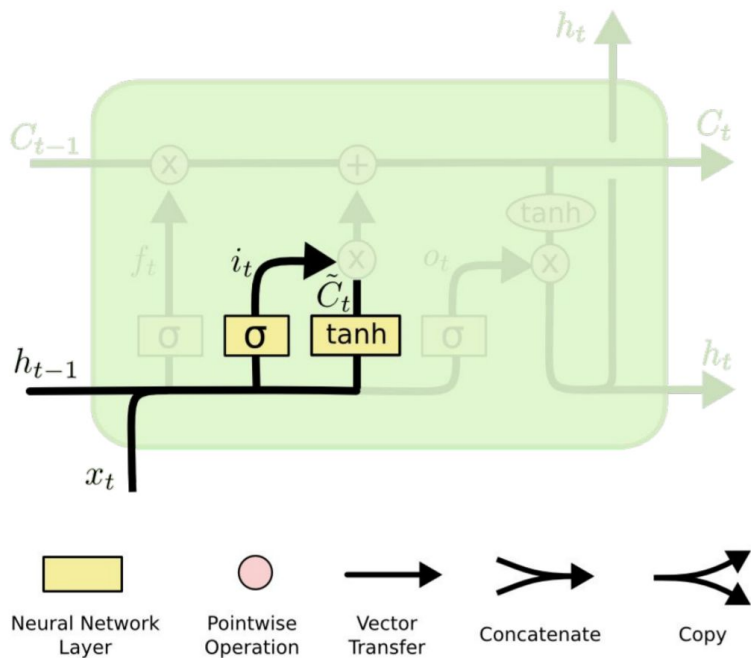
LSTM



Forget Gate:

$$f_t = \sigma (W_f \cdot \underbrace{[h_{t-1}, x_t]}_{\text{Concatenate}} + b_f)$$

LSTM



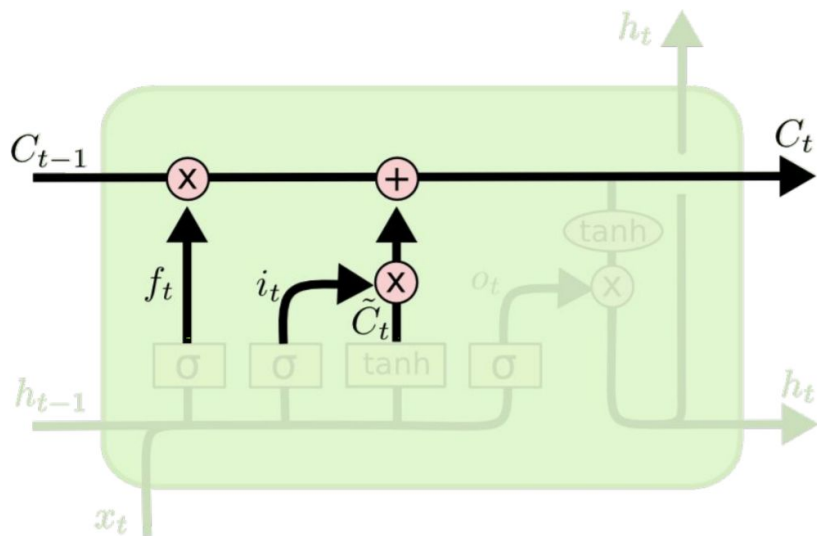
Input Gate Layer

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

New contribution to cell state

$$\tilde{C}_t = \underbrace{\tanh(W_C \cdot [h_{t-1}, x_t] + b_C)}_{\text{Classic neuron}}$$

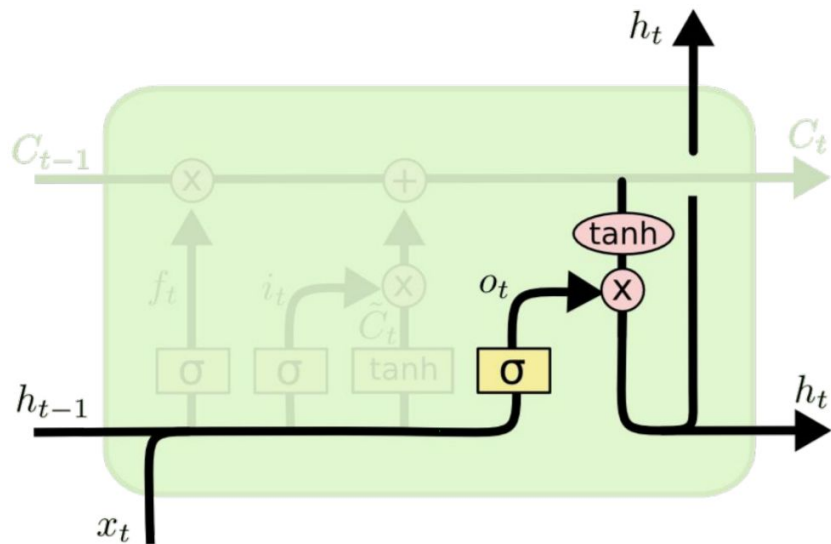
LSTM



Update Cell State (memory):

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

LSTM



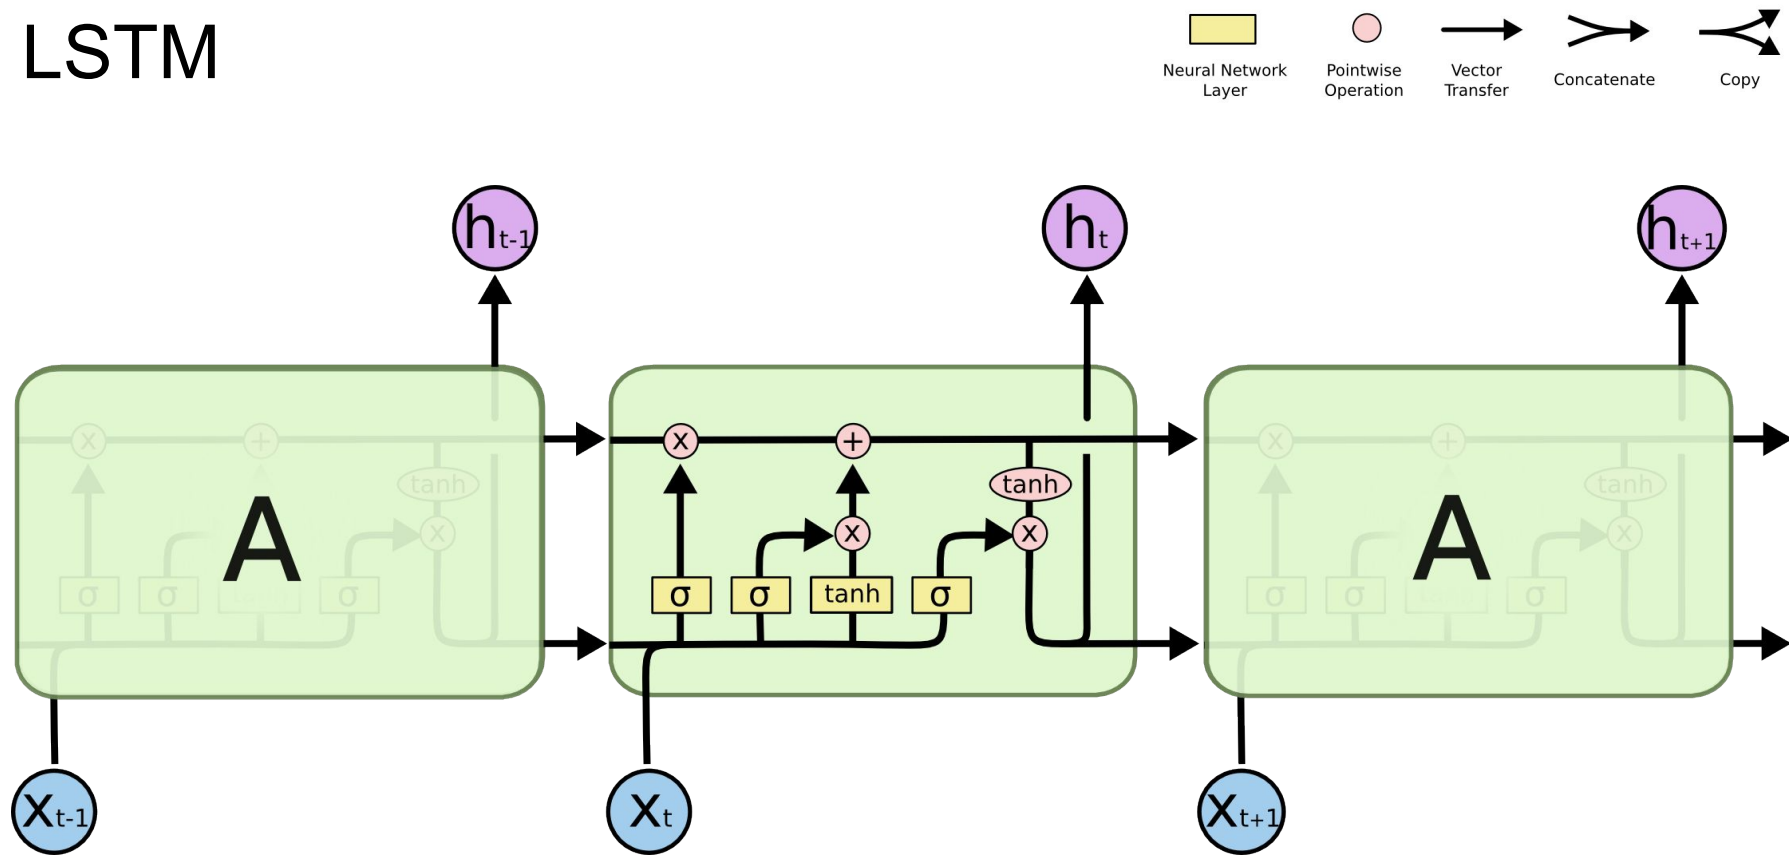
Output Gate Layer

$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

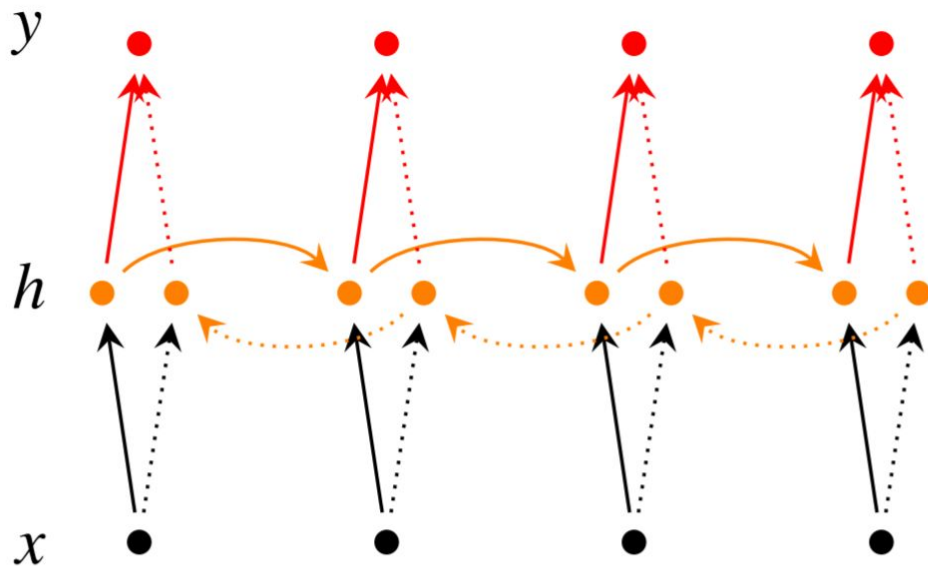
Output to next layer

$$h_t = o_t * \tanh (C_t)$$

LSTM



Bidirectional RNNs



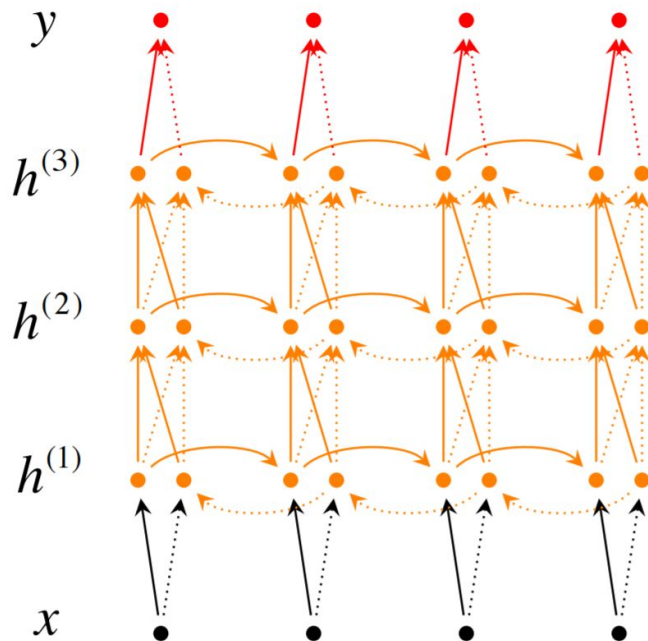
$$\vec{h}_t = f(\vec{W}x_t + \vec{V}\vec{h}_{t-1} + \vec{b})$$

$$\overleftarrow{h}_t = f(\overleftarrow{W}x_t + \overleftarrow{V}\overleftarrow{h}_{t+1} + \overleftarrow{b})$$

$$y_t = g(U[\vec{h}_t; \overleftarrow{h}_t] + c)$$

$h = [\vec{h}; \overleftarrow{h}]$ now represents (summarizes) the past and future around a single token.

Stacking RNNs



$$\vec{h}_t^{(i)} = f(\vec{W}^{(i)} h_t^{(i-1)} + \vec{V}^{(i)} \vec{h}_{t-1}^{(i)} + \vec{b}^{(i)})$$

$$\overleftarrow{h}_t^{(i)} = f(\overleftarrow{W}^{(i)} h_t^{(i-1)} + \overleftarrow{V}^{(i)} \overleftarrow{h}_{t+1}^{(i)} + \overleftarrow{b}^{(i)})$$

$$y_t = g(U[\vec{h}_t^{(L)}; \overleftarrow{h}_t^{(L)}] + c)$$

Each memory layer passes an intermediate sequential representation to the next.

References

1. Stanford NLP course : <http://web.stanford.edu/class/cs224n/syllabus.html>
2. Wild ML blog post :
<http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>
3. Colah's blog post : <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
4. Bengio's slides :
<https://drive.google.com/open?id=0ByUKRdiCDK7-LXZkM3hVSzFGTkE>
- 5.